

THE COMPOUNDING BOTTLENECK: A UNIFIED CAPACITY ACCOUNT OF RETRIEVE-THEN-COMPRESS RETRIEVAL-AUGMENTED GENERATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Efficient retrieval-augmented generation (RAG) funds two fixed-dimensional vector bottlenecks out of one budget: the dense-retrieval embedding of dimension d_r , and the soft-context compression code of size $m \cdot d_c$. Prior work bounds the retrieval stage or studies compression overflow, but never the composition. We give a unified capacity account that treats both stages as thresholded-separation problems of the same geometric form, and we make four observations, each validated from controlled free-vector geometry up to production encoders. First, the two stages carry sharp capacity walls whose critical dimension grows with content; the retrieval wall replicates across two embedder families, two datasets, and optimal (PCA) truncation, and on the adversarial LIMIT benchmark three real embedders collapse to 2.8–8.2% recall@100. Second, composing the stages *compounds*: the best split of a shared budget underperforms either stage given the full budget, and the sign of the deviation from the product law is governed by the correlation ρ between which items are hard to retrieve and which are hard to compress. Third, we measure ρ on a real pipeline and find $\rho \approx 0$, so real RAG sits in the multiplicative regime and a simple product-law allocation is near-exact; the allocation is interior, which we confirm both on a synthetic corpus and on a real multi-hop QA benchmark with a frozen reader scored by EM/F1 (a single-peaked allocation curve and a 0.048 F1 compounding gap at a tight budget, over three seeds), and deployed big-embedder/tiny-compressor configurations sit far off the frontier. Fourth, a multi-view “facet-lens” escapes the wall in the free-vector worst case, but we show honestly that the escape does not transfer to strong real encoders, which redirects the practical lever back to allocation. All claims are reproducible from committed configurations.

1 INTRODUCTION

Retrieval-augmented generation has settled into a common efficient template: a query is embedded and matched against a corpus by a dense retriever (Karpukhin et al., 2020; Lewis et al., 2020), the top passages are compressed into a handful of soft tokens so the generator pays little attention cost (Mu et al., 2023; Ge et al., 2023; Cheng et al., 2024), and a frozen decoder answers from those tokens. Both stages are, at heart, a fixed-dimensional vector bottleneck: retrieval squeezes relevance through an embedding of dimension d_r , and compression squeezes a passage through a code of size $m \cdot d_c$. The two bottlenecks are usually studied apart. On the retrieval side, a line of work now argues that a d_r -dimensional single-vector index cannot realize arbitrary top- k relevance patterns, and Weller et al. (2025) make this precise by tying feasibility to the sign-rank of the relevance matrix and by building the LIMIT stress test on which frontier models fail. On the compression side, Belikova et al. (2026) observe empirically that soft tokens “overflow” past a finite capacity, and Liu et al. (2026) argue on information-theoretic grounds that full compression is neither feasible nor necessary, yet decline to give a capacity bound or to connect compression back to the embedding-dimension wall.

What has been missing is the composition. Efficient RAG runs the compressor on the *output* of the retriever, so the two capacity-limited stages act in series, and the natural question is whether their




Figure 1: **The compounding bottleneck.** (a) Efficient RAG passes a query through two fixed-dimensional vector bottlenecks in series—the retrieval embedding (d_r) and the soft compression code ($m \cdot d_c$)—so the realizable set is the intersection of two subspace-constrained families and $c_{\text{eff}} \leq \min(c_{\text{retrieve}}, c_{\text{compress}})$, strict under generic misalignment. (b) Composing two stages that each achieve recall $p < 1$ yields p_{RPC} under independence, which lies below either stage alone; the shaded region is the compounding gap.

errors simply add, or whether stacking two bottlenecks does something worse. We find that they *compound*: there are inputs on which retrieval alone succeeds and on which compressing the gold passage alone succeeds, and yet the composed pipeline fails, and the size of this effect is predictable from the geometry of the two representations rather than from either stage’s standalone metric. This breaks the field’s implicit assumption that retrieval error and compression error are independent and additive, and it has a direct practical consequence, since a budget spent overwhelmingly on a large embedding while the compressor is starved sits far from the optimal allocation.

We organize the paper around four contributions, illustrated in Figure 1. (C1) We cast both stages as thresholded-separation problems of the same geometric form and show each has a sharp capacity wall whose critical dimension grows with content; the retrieval wall is not an artifact of one embedder, one dataset, or one truncation scheme, and it appears in full force on Weller’s own adversarial data. (C2) We show that composing two sub-full stages compounds, we characterize the deviation from the product law through a hardness-correlation ρ , and—going beyond the prior framing that left ρ hypothetical—we *measure* it on a real pipeline and find it near zero, which places real RAG squarely in the multiplicative regime. (C3) Under a fixed shared budget we derive and validate the optimal split, on synthetic corpora and end-to-end on real encoders, and show that deployed configurations are off the frontier. (C4) We test the obvious escape—a multi-view facet-lens that gives each semantic facet its own low-rank subspace—and report the honest result that it escapes the wall in the free-vector worst case but *not* on strong real encoders, which we explain and which sharpens rather than weakens the account. Every headline number is validated twice, from a free-vector best case that upper-bounds any encoder to production embedders, and the negatives are reported alongside the positives because they are what map the boundary of the effect.

2 RELATED WORK

Limits of dense retrieval. That a fixed embedding dimension caps what a single-vector retriever can express is the premise we build on. Weller et al. (2025) connect the number of realizable top- k document subsets to embedding dimension through sign-rank (Razborov & Sherstov, 2010; Forster, 2002) and release LIMIT, where even strong models score low; scaling studies of embedding dimension in IR and analyses of granularity failures point the same way. Matryoshka representation learning (Kusupati et al., 2022) makes the leading coordinates of an embedding independently useful, which is exactly what lets us trace the wall by truncation on a real encoder. Multi-vector late interaction (Khattab & Zaharia, 2020; Santhanam et al., 2022) is the known partial escape from the single-vector bound, and we position our facet-lens against it. Standard dense retrievers and embedders we use or reference include Reimers & Gurevych (2019); Karpukhin et al. (2020); Izacard et al. (2022); Wang et al. (2022); Li et al. (2023b); Xiao et al. (2023); Merrick et al. (2024); Ni et al. (2022), benchmarked in the style of Thakur et al. (2021); Muennighoff et al. (2023).

Context and soft-prompt compression. Soft-compression methods replace a passage with a few learned vectors read by a frozen decoder: Gist tokens (Mu et al., 2023), in-context autoencoding (Ge et al., 2023), context-adapted models (Chevalier et al., 2023), and one-token extreme compression (Cheng et al., 2024), alongside hard-prompt pruning (Jiang et al., 2023a;b; Pan et al., 2024; Li et al., 2023a) and abstractive compression (Xu et al., 2024; Li et al., 2024). Closest to us, Belikova et al. (2026) detect an “overflow” regime where compressed tokens can no longer answer a query, and Liu et al. (2026) argue that query-conditioned selection beats full compression; both stop short of a capacity law and of any link to the retrieval wall. We supply that link and, crucially, the cross-stage composition.

RAG systems. The broader pipeline we abstract is instantiated by Lewis et al. (2020); Guu et al. (2020); Borgeaud et al. (2022); Izacard et al. (2023); Izacard & Grave (2021); Ram et al. (2023); Shi et al. (2024); Asai et al. (2024). Our contribution is orthogonal: not a better pipeline, but a capacity account of why stacking two fixed-dimensional codes loses information that either stage alone would keep, and what to do about it.

3 SETUP AND THE CAPACITY PRIMITIVE

Notation. A corpus is $D = \{d_1, \dots, d_N\}$ with queries q and a ground-truth relevance matrix $A \in \{0, 1\}^{|Q| \times N}$, where row i marks the documents relevant to query i . The retrieval stage uses a document encoder $\varphi : d \mapsto x \in \mathbb{R}^{d_r}$ and a query encoder $\psi : q \mapsto y \in \mathbb{R}^{d_r}$, scores by $s(q, d) = \langle y, x \rangle$, and returns the top k . Realizing A by thresholding $YX^\top$ is governed by the sign-rank of $2A - 1$ (Weller et al., 2025), which is the retrieval wall. The compression stage maps a passage P to $Z \in \mathbb{R}^{m \times d_c}$ (m soft tokens of width d_c) that a frozen reader answers from; we cast “can the reader recover the queried fact from Z among distractors” as a second thresholded-separation problem with its own capacity c_{compress} . Because the compressor acts on the retriever’s output, the realizable set of the pipeline is the intersection of two subspace-constrained families, so $c_{\text{eff}} \leq \min(c_{\text{retrieve}}, c_{\text{compress}})$, with strict inequality generically, since the retriever’s discriminative subspace and the compressor’s reconstructive subspace need not align.

The free-embedding capacity measure. To measure capacity without confounding it with a particular encoder’s quality, we directly optimize unconstrained document and query vectors of a given dimension to realize a target pattern, using a pairwise margin objective, and report the fraction of queries for which every relevant document outranks every non-relevant one (“realizability”). This is the best case for *any* encoder of that dimension, so it upper-bounds real models and isolates the geometric limit; where we later want a smoother signal we report $\text{recall}@k$ instead. The same primitive drives the retrieval probe (optimize $\mathbb{R}^{d_r}$ vectors against a relevance pattern) and the compression probe (a slot-memory read head over m codes of width d_c), which is what lets us compare the two walls on the same axis. All runs are seeded and multi-seed for headline claims, and we report compute alongside accuracy; experiments ran on single commodity GPUs (Kaggle T4/P100).

4 TWO WALLS OF THE SAME GEOMETRY (C1)

The retrieval wall. On the adversarial “all-pairs” pattern, where every document pair is some query’s gold set ($k = 2$), free-embedding realizability shows a sharp phase transition in d_r : for a corpus of $n_d = 40$ it jumps from 0.38 to 1.0 between $d_r = 4$ and $d_r = 8$ across three seeds with near-zero variance, and the critical dimension d_r^* moves outward as the corpus grows. The growth is real but slow: on a five-seed, two-pattern sweep (all-pairs and random 3-subsets) d_r^* moves only from 8 to 11 as n_d goes from 40 to 400, consistent with a sub-logarithmic law (a power fit gives exponent $b \approx 0.16 - 0.18$), and we note honestly that a tidy “ $d_r^* \propto d/n_d$ ” sub-critical collapse we initially conjectured does *not* hold ($R^2 = 0.18$); what survives, and is all the allocation law needs, is that a modest critical dimension exists and grows with content.

The wall is not a free-vector artifact. Truncating a Matryoshka encoder (Kusupati et al., 2022) and re-normalizing traces the same ramp-then-saturate curve on real data: `mx-bai-embed-large` on SciFact climbs $\text{recall}@10$ from 0.09 to 0.87 and is statistically flat beyond $d_r = 512$, so the top half of a 1024-dimensional production embedding is wasted budget. Table 1 shows this replicates across

Table 1: **The retrieval wall generalizes** across a second embedder, a second dataset, and optimal (PCA) truncation. We report recall@10 at truncated dimension d_r and the fraction of full-dimension recall retained at $d_r = 256$. In every case a modest d_r captures the bulk and the remaining dimensions add little.

Embedder / data (trunc.)	full d	$d_r=16$	32	64	128	256	full	@256/full
mxbai / SciFact (Matryoshka)	1024	0.33	0.57	—	0.79	0.83	0.87	95%
arctic-m / SciFact (Matryoshka)	768	0.18	0.39	0.59	0.75	0.82	0.85	97%
mxbai / FiQA (Matryoshka)	1024	0.10	0.24	0.42	0.53	0.60	0.65	92%
bge-large / SciFact (PCA)	1024	0.49	0.65	0.76	0.80	0.85	0.86	98.5%

Table 2: **The wall on canonical adversarial data** (LIMIT, Weller et al., 2025). Full set: 50k documents, 1000 queries, two gold per query. All three real embedder families collapse; the failure is a property of the single-vector paradigm, not a model quirk.

Embedder	LIMIT-full (50k docs)				LIMIT-small (46)
	R@2	R@10	R@100	R@1000	R@10
mxbai-embed-large	0.004	0.010	0.028	0.103	0.441
arctic-embed-m	0.009	0.024	0.082	0.254	0.619
bge-base	0.007	0.015	0.045	0.163	0.492

settings that rule out the obvious objections: a *second* embedder family (arctic-embed-m, 768d) reaches 97% of full recall by $d_r = 256$; a *second* dataset (FiQA, a domain shift) reaches 92% by $d_r = 256$; and under *PCA* truncation—the optimal d_r -dimensional linear subspace, not a Matryoshka ordering—bge-large saturates hard at $d_r = 256$ (98.5% of full, the top 768 dimensions adding 0.013), which is the cleanest wall of the set and the one that kills the “your wall is just Matryoshka” reading.

Most tellingly, the wall appears on the field’s own adversarial benchmark. On the full LIMIT set (Weller et al., 2025) (50k documents, two gold per query) three real embedders collapse to recall@100 of 0.028, 0.082, and 0.045 (Table 2); on the 46-document LIMIT-small, recall@10 is only 0.44–0.62, i.e. returning ten of just forty-six documents still misses half the gold. That three different embedder families fail together, and consistently, is what makes this the wall of the single-vector *paradigm* rather than any one model’s quirk.

The compression wall. The compression stage carries a wall of the same form. On controlled passages of n_f facts read from m soft codes, slot-memory realizability transitions at a critical code size $D_c^* = m \cdot d_c \approx 2n_f$, growing roughly linearly with content, and the transition is softer than the retrieval one but unmistakable. The wall survives on a real encoder: a single 1024-dimensional mxbai embedding used as the compressed code holds about one fact, with a light probe recovering the queried value at 0.98, 0.53, 0.32, 0.24, 0.20 for $n_f = 1, 2, 4, 8, 16$ (chance 0.125), and giving each fact its own slot restores capacity to 0.96–0.99 once the per-slot width clears a floor. We also flag a scope boundary honestly: under a learned attention-read code, capacity is not shape-invariant—one fat token and many thin tokens both collapse while an interior slot count is best—but under real-encoder truncation only the thin-slot collapse remains, so the interior optimum is specific to the learned-code mechanism rather than a universal law.

5 COMPOSING THE TWO STAGES COMPOUNDS (C2)

Independent composition already bites. If the two stages fail on independent items, the pipeline recall is the product $p_R(d_r) \cdot p_C(D_c)$, so composing two sub-full stages must underperform either one at the full budget. On a synthetic corpus that shares a budget $B = d_r + D_c$ between the real retrieval and compression walls of Section 4, the best split at $B = 64$ reaches only 0.559 even though each stage alone exceeds 0.95 at that budget—a compounding gap of 0.389—and the gap closes to 0.001 by $B = 192$ once both stages clear their walls. The effect is largest exactly where deployments live, at tight budgets where neither stage is saturated.




Figure 2: **C1**. Left: the free-embedding retrieval wall—realizability jumps sharply once d_r clears a critical value that grows with corpus size n_d . Right: on LIMIT, three real embedder families all fail the adversarial pattern at full dimension.




Figure 3: **C2–C3 on real models**. Left: end-to-end accuracy versus the retrieval share of a fixed budget—the optimum is interior, and the big-embedder/tiny-compressor corner is off it. Right: measured $\rho \approx 0$; the copula (product law) reproduces the observed pipeline recall at every operating point.

The deviation from the product law is set by a hardness correlation. Independence is a baseline, not a law, so we model the two per-item success events with a Gaussian copula that holds the marginals fixed at the measured walls and varies the correlation ρ between which items are hard to retrieve and which are hard to compress. The limits are exact: $\rho = 0$ recovers the product p_{RPC} ; aligned hardness ($\rho > 0$) makes failures redundant and lifts the pipeline toward $\min(p_R, p_C)$; and misaligned hardness ($\rho < 0$) is the only regime that compounds *super*-multiplicatively, pushing the pipeline below p_{RPC} toward the Fréchet floor $\max(0, p_R + p_C - 1)$. Concretely, at a balanced 32:32 split the pipeline moves from 0.515 at $\rho = -0.5$ (below the independent 0.559) up to 0.687 at $\rho = +0.9$. So the sign of the compounding is an empirical question about real corpora, which prior work left open.

We measure ρ , and it is near zero. On a real pipeline—a truncated mxbai retriever and a real-encoder compressor over one shared corpus—we record, per query, a retrieval difficulty (gold-versus-best-distractor margin) and a compression difficulty (the reader’s correct-versus-best-other logit margin), and correlate them. Across three balanced, sub-saturated operating points the mean ρ_ϕ is $+0.009$, i.e. essentially zero, and the copula prediction reproduces the observed pipeline recall almost exactly: observed $\{0.234, 0.443, 0.555\}$ against the product law $\{0.227, 0.439, 0.555\}$, a mean absolute error of 0.0035 (independent) or 0.0012 (with the measured ρ). The reading is that real RAG, at least on a benign corpus, lives in the multiplicative regime, so the pessimistic super-multiplicative case needs adversarial anti-correlation that we do not see by default—which is reassuring for allocation, since it means p_{RPC} is a near-exact objective to optimize.




Figure 4: **C2–C3 on real multi-hop QA** (HotpotQA, Qwen2.5-1.5B reader, EM/F1, three seeds). End-to-end F1 versus the retrieval share of a fixed budget. The curve is single-peaked with an interior optimum (red), and past it more retrieval budget hurts (it starves compression). The shaded band is the compounding gap—the best split underperforms the weaker standalone stage—which is positive at the tight budget $B=128$ (0.048 ± 0.007 F1) and vanishes by $B=256$.

6 BUDGET-OPTIMAL ALLOCATION (C3)

Because the pipeline is (near-)multiplicative, the optimal split of a fixed budget follows a simple recipe: fund retrieval up to its wall d_r^* , then pour the remainder into compression, since dimensions past the retrieval wall buy almost nothing while the compressor is still starved. On the synthetic composition this is exactly what we see—the optimal split is $48:(B-48)$ for all $B \geq 96$, with a balanced 32:32 preferred only when B is below twice the wall—and it holds end-to-end on real encoders. On one shared corpus with a real mxbai retriever and a real-encoder compressor, the best split at $B = 128$ is interior (64:64) and reaches 0.234 while the stages alone reach 0.63 and 0.89, a compounding gap of 0.392 that shrinks to 0.137 by $B = 256$ (Figure 3, left). The practical statement is blunt: the standard configuration of a large $d_r \approx 768-1024$ embedding with a tiny $m \approx 1-8$ compressor sits far off the frontier, since it spends heavily on embedding dimensions past the retrieval wall while the compressor overflows, and re-allocating the same budget buys accuracy for free. Combined with C1, which locates each wall cheaply by truncation, this is an actionable procedure: measure both walls, then split the budget so retrieval is funded to d_r^* and the rest goes to compression.

Confirmation on a real multi-hop QA benchmark (E9). The synthetic pipeline of E6 isolates the mechanism, but the claim that matters is whether it survives on an actual question-answering task with a real frozen reader. We pool every HotpotQA question’s paragraphs into one shared corpus (a genuine retrieval problem, not one isolated per question), truncate a real embedder to d_r for retrieval, select sentences into a fixed token budget by D_c -truncated query similarity (the deployable, selective instantiation of compression used by RECOMP-style systems), and let a frozen Qwen2.5-1.5B-Instruct answer, scored by EM/F1 over three seeds. Both claims survive on real answer quality (Figure 4). The allocation curve is single-peaked with an *interior* optimum: F1 rises to a peak near $d_r \approx 80-96$ and then declines along a long tail out to $d_r=224$, so past the optimum more retrieval budget *hurts*, because it starves compression— exactly the shape C3 predicts. And composition compounds at the tight budget: at $B = 128$ the best split reaches F1=0.335 against standalone references of 0.39 (retrieval) and 0.56 (compression), a compounding gap of 0.048 ± 0.007 F1 with all three seeds positive; the gap shrinks to 0.009 ± 0.015 (null) by $B = 256$ once both stages clear their walls, the same shrinking-with-budget pattern as the synthetic result and consistent with the measured $\rho \approx 0$. The precise optimum location is noisy across seeds (d_r^* in $64-160$), but its interior character and the tight-budget compounding gap are robust—the real-QA, EM/F1 counterpart of E3 and E6.




Figure 5: **C4 and its scope.** Left (free vectors): routed facet-lenses escape the wall at half the budget of a single vector, while generic multiview is worse than single. Right (frozen real encoder): the escape vanishes—a single learned projection ties the routed lenses and leads at tight budgets—because a strong encoder has already linearized the facets.

7 WHEN DOES MULTI-VIEW HELP? THE ESCAPE AND ITS SCOPE (C4)

A natural way to beat a single-vector wall is to stop using a single vector. We test a facet-lens: K low-rank lenses, each specialized to one semantic facet (entities, numerics, relations, negation, ...), with a lightweight router sending each query to its facet’s lens, so every doc gets K small codes instead of one large one at equal total budget. In the free-vector worst case this escapes the wall cleanly. On a facet-structured pattern, routed facet-lenses realize the relevance structure at critical budget $d_r^* = 12$ against a single vector’s 24—half the budget—and, tellingly, a generic K -view MaxSim scheme with no routing is *worse* than a single vector ($d_r^* = 48$), so the lever is specialization plus routing, not multi-vector per se. This distinguishes the facet-lens from late-interaction multi-vector retrieval (Khattab & Zaharia, 2020), where the gain comes from many token vectors rather than from routing to a facet.

We then report the honest negative: the escape does *not* transfer to a strong real encoder. Learning low-rank lenses over frozen mx-bai embeddings of a facet-structured corpus (ten facets), a single *learned* projection ties the routed lenses—both reach mean average precision 0.9 at budget 40—and the single projection actually leads at tight budgets (0.735 versus 0.581 at budget 20), while generic multiview is far behind (budget 320). The mechanism is instructive: a powerful pretrained encoder has already spent representational capacity making the facets close to linearly separable, so a single learned projection can unmix them and routing adds little, whereas the free-vector escape was a worst-case sign-rank phenomenon in which no single set of vectors could satisfy all facet constraints at once. The consequence for practice is clean and, we think, more useful than a method that only wins in the worst case: on real encoders the lever is allocation (C3), not multi-view, and multi-view is not a free lunch. We therefore keep C4 as a scoped “when does it help” result rather than the paper’s headline, which protects the empirical core (C1–C3) from an overclaim.

8 DISCUSSION AND LIMITATIONS

The account is deliberately a capacity characterization rather than a closed-form theorem. The free-vector measure is a best case, so real encoders can only do worse, which is why the wall shows up so cleanly on truncated production embeddings and on LIMIT; but it also means our $c_{\text{eff}} \leq \min(\cdot)$ statement is a proposition about the realizable set backed by a direct empirical measurement of that set, not a bound with constants. Several boundaries are worth stating plainly. The compounding experiments use controlled single-task corpora so that we can hold the marginals fixed and vary ρ ; the measured $\rho \approx 0$ is on a benign corpus, and whether adversarial inputs can drive the *joint* ρ negative—LIMIT shows retrieval can be driven near zero, but the joint hardness correlation under such pressure is untested—is the one genuinely open question we would flag, since a negative ρ is where the pessimistic super-multiplicative regime lives. The compression capacity is isolated with a light probe rather than a fully-trained generative soft-token compressor, which in our hands stays at chance without a large training budget; the probe cleanly establishes the capacity ceiling but does

not model the end-to-end training dynamics of methods like Cheng et al. (2024); Ge et al. (2023). Finally, C4’s escape is scoped out of the strong-real-encoder regime, so we do not claim multi-view as a general remedy; we claim, more narrowly and we think more usefully, that specialization-plus-routing is what would be needed for an escape and that a strong encoder makes it redundant.

9 CONCLUSION

We treated the retriever and the compressor of efficient RAG as two fixed-dimensional vector bottlenecks of the same geometric form, and showed that composing them compounds: each stage has a capacity wall that grows with content and that appears on real encoders, on a second dataset, under optimal truncation, and on the adversarial LIMIT set; the composition underperforms either stage, with the deviation from the product law set by a hardness correlation that we measured to be near zero, which places real RAG in a multiplicative regime; and the resulting allocation is interior, so the deployed habit of a large embedding with a tiny compressor is off the frontier and can be improved at equal cost. We also tested the obvious multi-view escape and reported honestly that it does not transfer to strong real encoders, which points the practical lever back at allocation. The picture that emerges is a full arc—a diagnosis of the problem, an actionable allocation for it, and a careful account of when a fancier representation would and would not help—and we hope it reframes “efficient RAG” as a budget-allocation problem across two coupled bottlenecks rather than two tricks tuned in isolation.

REPRODUCIBILITY STATEMENT

Every experiment is driven by a committed YAML configuration and a fixed seed, headline claims use at least three seeds with reported variance, and each run emits a machine-readable results file and a human-readable summary. The capacity primitives, synthetic-corpus generators, real-encoder truncation and probe code, the copula composition, and the facet-lens are all in the released source, and the figures and tables here are produced directly from the committed outputs. Real-model experiments use publicly available encoders (mxbai-embed-large-v1, arctic-embed-m-v1.5, bge-large-en-v1.5, all-MiniLM-L6-v2) and datasets (SciFact, FiQA, LIMIT) and run on single commodity GPUs.

REFERENCES

- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.11511.
- Julia Belikova, Danila Rozhevskii, Dennis Svirin, Konstantin Polev, and Alexander Panchenko. Detecting overflow in compressed token representations for retrieval-augmented generation. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Student Research Workshop)*, 2026. arXiv:2602.12235.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, et al. Improving language models by retrieving from trillions of tokens. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, 2022.
- Xin Cheng, Xun Wang, Xingxing Zhang, Tao Ge, Si-Qing Chen, Furu Wei, Huishuai Zhang, and Dongyan Zhao. xRAG: Extreme context compression for retrieval-augmented generation with one token. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2405.13792.
- Alexis Chevalier, Alexander Wettig, Anirudh Ajith, and Danqi Chen. Adapting language models to compress contexts. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- Jürgen Forster. A linear lower bound on the unbounded error probabilistic communication complexity. *Journal of Computer and System Sciences*, 65(4):612–625, 2002.

- Tao Ge, Jing Hu, Lei Wang, Xun Wang, Si-Qing Chen, and Furu Wei. In-context autoencoder for context compression in a large language model. *arXiv preprint arXiv:2307.06945*, 2023.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. REALM: Retrieval-augmented language model pre-training. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020.
- Gautier Izacard and Edouard Grave. Leveraging passage retrieval with generative models for open domain question answering. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2021.
- Gautier Izacard, Mathilde Caron, Lucas Hosseini, Sebastian Riedel, Piotr Bojanowski, Armand Joulin, and Edouard Grave. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research (TMLR)*, 2022. arXiv:2112.09118.
- Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. Atlas: Few-shot learning with retrieval augmented language models. *Journal of Machine Learning Research (JMLR)*, 24: 1–43, 2023.
- Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. LLMingua: Compressing prompts for accelerated inference of large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023a.
- Huiqiang Jiang, Qianhui Wu, Xufang Luo, Dongsheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. LongLLMingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression. *arXiv preprint arXiv:2310.06839*, 2023b.
- Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 6769–6781, 2020. arXiv:2004.04906.
- Omar Khattab and Matei Zaharia. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 39–48, 2020.
- Aditya Kusupati, Gantavya Bhatt, Aniket Rege, Matthew Wallingford, Aditya Sinha, Vivek Ramanujan, William Howard-Snyder, Kaifeng Chen, Sham Kakade, Prateek Jain, and Ali Farhadi. Matryoshka representation learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pp. 30233–30249, 2022.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv:2005.11401.
- Yucheng Li, Bo Dong, Frank Guerin, and Chenghua Lin. Compressing context to enhance inference efficiency of large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023a.
- Zehan Li, Xin Zhang, Yanzhao Zhang, Dingkun Long, Pengjun Xie, and Meishan Zhang. Towards general text embeddings with multi-stage contrastive learning. *arXiv preprint arXiv:2308.03281*, 2023b.
- Zongqian Li, Yidong Su, and Nigel Collier. 500xCompressor: Generalized prompt compression for large language models. *arXiv preprint arXiv:2408.03094*, 2024.
- Yunhao Liu, Zian Jia, Xinyu Gao, Kanjun Xu, and Yun Xiong. Rethinking soft compression in retrieval-augmented generation: A query-conditioned selector perspective. In *Proceedings of the ACM Web Conference (WWW)*, 2026. arXiv:2602.15856.
- Luke Merrick, Danmei Xu, Gaurav Nuti, and Daniel Campos. Arctic-embed: Scalable, efficient, and accurate text embedding models. *arXiv preprint arXiv:2405.05374*, 2024.

- Jesse Mu, Xiang Lisa Li, and Noah Goodman. Learning to compress prompts with gist tokens. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. MTEB: Massive text embedding benchmark. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2023. arXiv:2210.07316.
- Jianmo Ni, Chen Qu, Jing Lu, Zhuyun Dai, Gustavo Hernández Ábrego, Ji Ma, Vincent Y. Zhao, Yi Luan, Keith B. Hall, Ming-Wei Chang, and Yinfei Yang. Large dual encoders are generalizable retrievers. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022. arXiv:2112.07899.
- Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, Menglin Xia, Xufang Luo, Jue Zhang, Qingwei Lin, Victor Rühle, Yuqing Yang, Chin-Yew Lin, H. Vicky Zhao, Lili Qiu, and Dongmei Zhang. LLMLingua-2: Data distillation for efficient and faithful task-agnostic prompt compression. In *Findings of the Association for Computational Linguistics (ACL Findings)*, 2024. arXiv:2403.12968.
- Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. In-context retrieval-augmented language models. *Transactions of the Association for Computational Linguistics (TACL)*, 11:1316–1331, 2023.
- Alexander A. Razborov and Alexander A. Sherstov. The sign-rank of AC^0 . *SIAM Journal on Computing*, 39(5):1833–1855, 2010.
- Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019. arXiv:1908.10084.
- Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. ColBERTv2: Effective and efficient retrieval via lightweight late interaction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2022. arXiv:2112.01488.
- Weijia Shi, Sewon Min, Michihiro Yasunaga, Minjoon Seo, Rich James, Mike Lewis, Luke Zettlemoyer, and Wen-tau Yih. REPLUG: Retrieval-augmented black-box language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2024. arXiv:2301.12652.
- Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2021. arXiv:2104.08663.
- Liang Wang, Nan Yang, Xiaolong Huang, Binxing Jiao, Linjun Yang, Daxin Jiang, Rangan Majumder, and Furu Wei. Text embeddings by weakly-supervised contrastive pre-training. *arXiv preprint arXiv:2212.03533*, 2022.
- Orion Weller, Michael Boratko, Iftekhhar Naim, and Jinhyuk Lee. On the theoretical limitations of embedding-based retrieval. *arXiv preprint arXiv:2508.21038*, 2025.
- Shitao Xiao, Zheng Liu, Peitian Zhang, Niklas Muennighoff, Defu Lian, and Jian-Yun Nie. C-Pack: Packed resources for general Chinese embeddings. *arXiv preprint arXiv:2309.07597*, 2023.
- Fangyuan Xu, Weijia Shi, and Eunsol Choi. RECOMP: Improving retrieval-augmented LMs with context compression and selective augmentation. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.04408.

A SUMMARY OF EXPERIMENTS

Table 3 collects every experiment behind the claims in the main text, with the setting and the headline number, so the results can be traced to a single row. Each corresponds to a committed configuration and a fixed seed; headline claims use at least three seeds.

Table 3: All experiments, their setting, and the headline result. “Free” denotes the free-embedding capacity primitive of Section 3; “real” denotes a frozen production encoder. Budgets are embedding/code dimensions.

ID	Kind	Setting	Headline result
E1/E1c	free	all-pairs & random- k patterns; sweep d_r, n_d ; 5 seeds	sharp wall; d_r^* grows $8 \rightarrow 11$ as n_d $40 \rightarrow 400$; “ d/n_d ” law refuted ($R^2=0.18$)
E2	free	slot-memory read; sweep D_c, n_f	compression wall $D_c^* \approx 2n_f$
E4	real	mxbai / SciFact; Matryoshka trunc.	recall@10 $0.09 \rightarrow 0.87$; flat past $d_r=512$
E4b	real	arctic-m / SciFact; Matryoshka	97% of full recall by $d_r=256$
E4c	real	mxbai / FiQA (domain shift)	92% of full recall by $d_r=256$
E4d	real	bge-large / SciFact; PCA trunc.	saturates at $d_r=256$ (98.5%)
E5/E5b	real	mxbai code; light read probe	one vector $\approx$ one fact ($0.98 \rightarrow 0.20$); multi-token recovers to 0.99
E8	real	LIMIT; 3 embedders; 50k / 46 docs	recall@100 = 2.8/8.2/4.5% (full)
E3	free	shared budget $B=d_r+D_c$	compounding gap 0.389 at $B=64$
E3b	free	Gaussian copula over stage successes	sign of gap set by ρ ; Fréchet-bounded
E7	real	mxbai retriever + compressor; per-item margins	$\rho_\phi = +0.009$; product law exact to ~ 0.001
E6	real	mxbai; one corpus; $N=2000$	gap 0.392 at $B=128$; optimum interior 64:64
E9	real	HotpotQA; Qwen2.5-1.5B; EM/F1; 3 seeds	interior optimum; gap 0.048 ± 0.007 F1 at $B=128$, null at 256
C4	free	facet pattern; single / multiview / routed	routed $d_r^*=12$ vs single 24; multi-view 48
C4b	real	frozen mxbai; 10 facets; learned lenses	single ties routed ($d_r^*=40$); no transfer

B IMPLEMENTATION DETAILS

Free-embedding capacity primitive. For a target relevance pattern $A \in \{0, 1\}^{M \times N}$ we optimize unconstrained document vectors $X \in \mathbb{R}^{N \times d}$ and query vectors $Q \in \mathbb{R}^{M \times d}$ with Adam on a pairwise margin loss that, per query, pushes the smallest relevant score above the largest non-relevant score by a margin; realizability is the fraction of queries for which this holds strictly. Because X, Q are free, this upper-bounds any encoder of dimension d , so a wall here is a wall for every encoder. The all-pairs pattern sets every document pair as one query’s gold set ($k=2$); random- k draws n_q distinct k -subsets.

Compression probe. A passage of n_f facts is written as m codes of width d_c ; a query attends over the m codes with a small read head and decodes the queried value with cross-entropy. On real encoders the codes are truncated frozen embeddings, and only the light read head is trained, which converges reliably and generalizes to held-out passages, so it measures the code’s capacity rather than the trainability of a full soft-token compressor (which, in our hands, stays at chance without a large training budget).

Copula composition and measured ρ . We place a Gaussian copula on the two binary per-item success events with the marginals fixed at the measured stage recalls and correlation ρ , and estimate pipeline recall by Monte-Carlo; the limits $\rho \in \{-1, 0, +1\}$ recover the Fréchet lower bound, the product, and the minimum. To measure ρ on a real pipeline we take the phi coefficient of the two binary successes and the Pearson correlation of two continuous difficulty margins (retrieval gold-minus-best-distractor; compression correct-minus-best-other logit).

Facet-lens. Each lens is a learned linear map $\mathbb{R}^D \rightarrow \mathbb{R}^{d/K}$ applied to frozen document and query embeddings, scored by $\langle W_q^f q, W_d^f d \rangle$; the router sends a query of facet f to lens f , so the doc-side budget d is split into K equal codes. Single is one full-width lens; generic multiview keeps K lenses

but scores by the max over views without routing. We train all conditions with the same objective and budget and report mean average precision.

Compute. All runs use a single commodity GPU (Kaggle T4 or P100). Real-encoder runs encode the corpus once and cache it; wall-clock ranges from seconds (LIMIT-small; the free-vector sweeps) to about twenty-six minutes for the 50k-document LIMIT-full sweep across three embedders.